


# RL lab4

Name: 陳以瑄 Student ID: 109705001

- Experimental Results**

Screenshot of Tensorboard training curve



Screenshot of testing results on TD3 (with the old reward)

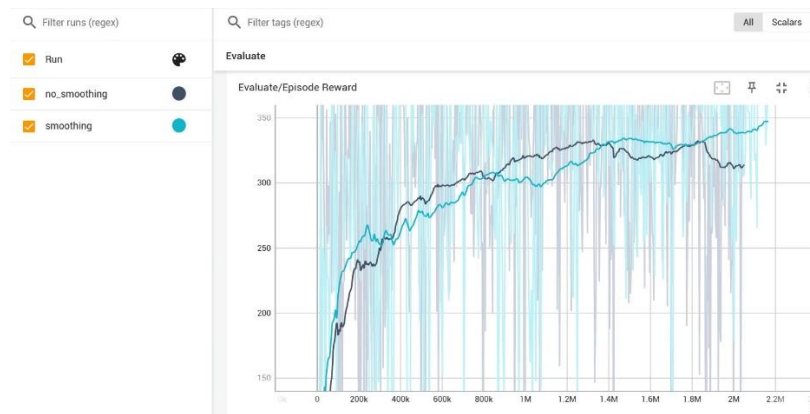
```
Evaluating...
Episode: 1    Length: 154    Total reward: 152.47
Episode: 2    Length: 331    Total reward: 385.72
Episode: 3    Length: 999    Total reward: 885.56
Episode: 4    Length: 292    Total reward: 325.77
Episode: 5    Length: 299    Total reward: 323.15
Episode: 6    Length: 144    Total reward: 117.64
Episode: 7    Length: 493    Total reward: 591.45
Episode: 8    Length: 342    Total reward: 386.23
Episode: 9    Length: 160    Total reward: 131.44
Episode: 10   Length: 999    Total reward: 871.33
average score: 417.0751429991862
```

Screenshot of testing results on TD3 (with the new reward)

```
Episode: 1    Length: 999    Total reward: 890.07
Episode: 2    Length: 984    Total reward: 901.50
Episode: 3    Length: 827    Total reward: 917.20
Episode: 4    Length: 999    Total reward: 893.44
Episode: 5    Length: 945    Total reward: 905.40
average score: 901.5217696232643
```

- Bonus2. Target Policy Smoothing**

Screenshot of Tensorboard training curve

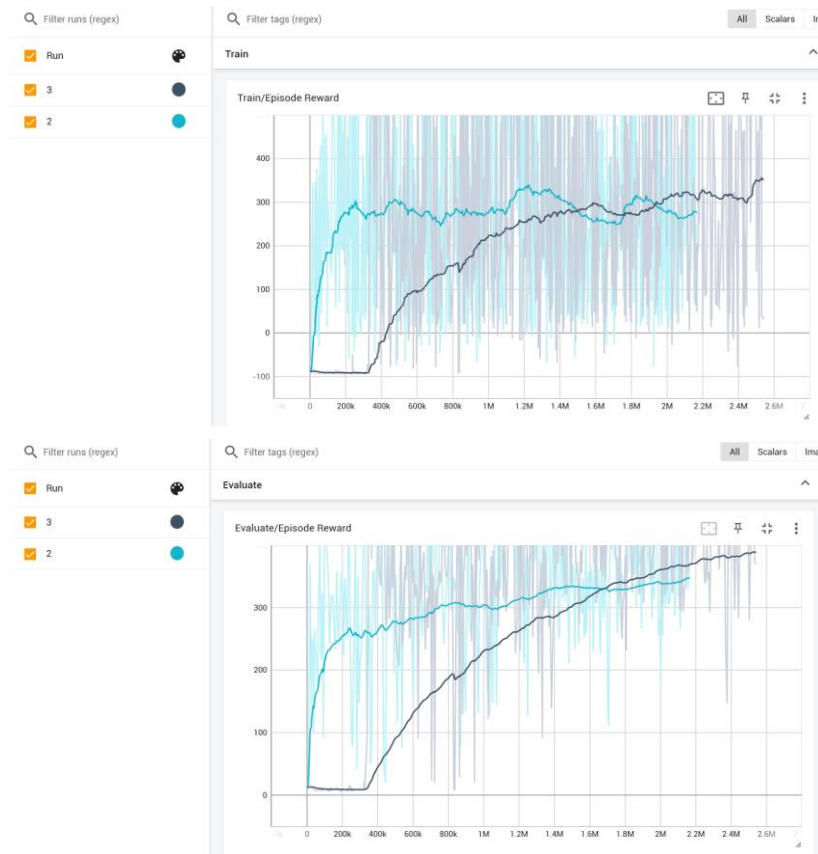


The plot above compares the results of the model with smoothing and without smoothing. As the figure shows, although they have only slight differences, we

can still observe that the model with smoothing performs better in the end. This is because smoothing makes the learning process less sensitive to small changes in the Q-values, resulting in greater stability.

- **Bonus3. Delayed Policy Update Mechanism**

Screenshot of Tensorboard training curve



I compared the model with a 'update\_freq' hyperparameter set to 3 to the one with a 'update\_freq' of 2. The results show that the model with a higher 'update\_freq,' which updates less frequently, learns more slowly in the beginning but performs better after 1.8 million updates. This is because delay update mechanism can stabilize the learning, so that it can get a better performance.

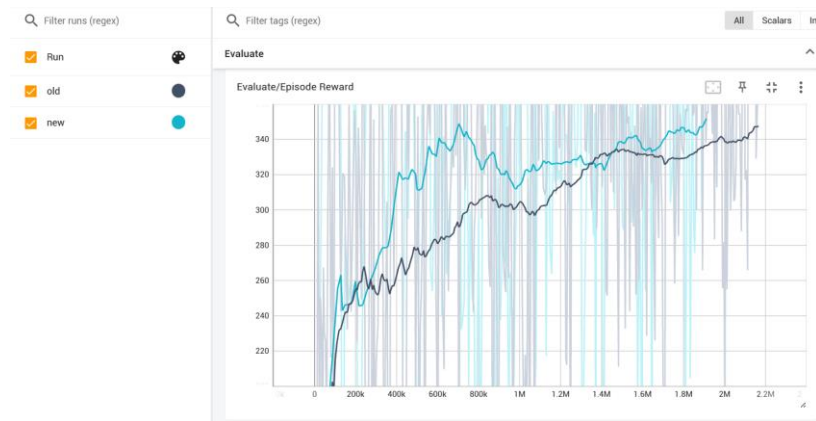
- **Bonus5. Reward Function Design**

The following is my reward function:

```
57 ———>#if terminates:
58 ———>#>reward = -150
59 ———># more road better
60 ———>reward += (float(road_pixel_count) - float(grass_pixel_count))/20.0
61
62 ———>if road_pixel_count < 10:
63 ———>terminates = True
64 ———>reward = -100
```

The idea is that if the proportion of the road surface in the visual scene is higher, it indicates that the race car is less deviated from the track. Therefore, it should receive a higher reward.

Screenshot of Tensorboard training curve



As the figure shows, the model with the new reward obtains higher rewards more quickly. Although its later improvement rate shows some stagnation, the overall performance still surpasses that of the old reward function. This means that the idea of encouraging the car to stay on the road can indeed help the model perform better.